

# Predictive Machine Learning System for Computer Price Estimation & Explainability

*A PREDICTIVE ML SYSTEM FOR ESTIMATING THE PRICE OF COMPUTER CONFIGURATIONS*

Group 6



# Team & Workflows



Data  
Preparation



Model  
Exploration



Additional  
models  
& Validation



Final model &  
Explainability



UI  
Development



NLP Integration  
& Interface



Documentation  
requirements

# Dataset Overview & Key Challenges

**8064×135 Spanish-language marketplace dataset**

**Numerical values embedded in text (e.g., '2,5 GHz', '1 TB')**

**Missing-Not-At-  
Random patterns:  
desktops lack  
batteries, budget  
laptops lack  
benchmarks**

**Redundant representations of same concept (screen size, CPU frequency)**

**Data raw to clean transformation**

Raw Field	Raw Example	Cleaned Value
15,6 pulgadas	15,6 pulgadas	15.6
1.024,00 € - 2.499,00 €	1.024,00 €	1024.00
512GB SSD NVMe	512GB SSD	512
Turbo máx. 4,3 GHz	4,3 GHz	4.3

**The dataset required not just cleaning, but conceptual understanding of hardware**

# Data Cleaning Strategy

**Built extraction functions to normalise frequencies, storage sizes, screen dimensions**

**Harmonised units (TB→GB, commas→decimal points, cm→inches)**

**Reduced from ~139 messy raw columns to ~44 structured hardware fields**

**Subsystem-based cleaning:**

- CPU
- GPU
- RAM
- Storage
- Display
- Physical
- Market fields

Stage	Columns	Notes
Raw dataset	135	Untidy text, duplicates
After cleaning	44	All numeric or categorical
After encoding	151	Model-ready numerical matrix
After feature selection	17	Only informative features

**The core of the project's accuracy comes from engineering high-quality, interpretable features which result in clear impact on the dataset as shown in the table**

# Data Cleaning & Missingness Patterns

== Missing % for ssd\_gb by category ==

Grouped by Tipo de producto:

Tipo de producto	
Kit ampliación PC	1.000000
Barebone	0.793651
Netbook	0.428571
Chromebook	0.381579
Ultrabook	0.375000
PC completo	0.333333
Portátil 3D	0.333333
Mini PC	0.270270
Portátil convertible	0.224719
Portátil profesional	0.222772
Portátil gaming	0.199229
Portátil multimedia	0.161841
PC de oficina	0.130856
PC multimedia	0.125714
Workstation	0.123188
PC gaming	0.044109
Thin Client	0.040000
Name: ssd_gb, dtype: float64	

Grouped by Tipo:

...	
Tipo	
Desktop	0.741071
Laptop	0.018105
Name: weight_kg, dtype: float64	

- Missing values follow structural patterns, not randomness
- Desktops show 74% missing weight — this is because sellers rarely include physical specs
- Barebones kits have 100% missing SSD capacity — storage is optional
- Laptops have very low weight missingness — buyers expect this info

Structural missingness (value doesn't exist by design):

- hdd\_gb: Missing = no HDD → fill with 0, add has\_hdd flag
- battery\_wh: Desktops have no battery → fill with 0, add has\_battery flag

Unknown values (data not recorded):

- height, weight\_kg: Median imputation by product type + missingness flags
- cpu\_mark, gpu\_mark: Model-based imputation using RandomForest
- Drop very sparse columns

# Missing Value Strategy

**Structural  
imputation:  
desktops have  
no battery →  
set to 0**

**Median  
imputation by  
device type for  
physical  
attributes**

**Model-based  
imputation  
using  
RandomForest  
to infer  
CPU/GPU  
marks where  
missing**

**Added  
missingness  
indicators  
("battery\_mis-  
sing",  
"height\_mis-  
sing")**

**Supports  
prediction  
when the user  
leaves fields  
blank –  
required by the  
assignment**

- CPU/GPU benchmarks are missing mostly in low-end devices → imputing them prevents price underestimation
- Battery\_wh missing always implies desktop, which we converted to 0
- Missing dimensions (height/weight) were filled using median per Tipo (Laptop/Desktop) to preserve category structure\

**Missingness carries semantic meaning; the model learns from it.**

# Feature Engineering & Encoding

```
Data columns (total 49 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   Título          7927 non-null    object  
 1   Precio_Rango     7927 non-null    object  
 2   Pantalla_Luminosidad 7927 non-null    object  
 3   Procesador_Procesador 7927 non-null    object  
 4   Gráfica_Tarjeta gráfica 7927 non-null    object  
 5   Tipo             7927 non-null    object  
 6   screen_size_inch 7927 non-null    float64 
 7   cpu_base_ghz     7927 non-null    float64 
 8   cpu_turbo_ghz   7927 non-null    float64 
 9   cpu_tdp_w        7927 non-null    float64 
 10  ram_gb          7927 non-null    float64 
 11  ssd_gb          7927 non-null    float64 
 12  hdd_gb          7927 non-null    float64 
 13  vram_gb         7927 non-null    float64 
 14  height           7927 non-null    float64 
 15  width            7927 non-null    float64 
 16  depth            7927 non-null    float64 
 17  weight_kg        7927 non-null    float64 
 18  chassis_color    7927 non-null    object  
 19  battery_wh       7927 non-null    float64 
 ...
 47  cpu_vendor       7927 non-null    object  
 48  gpu_vendor       7927 non-null    object  
dtypes: float64(19), int64(7), object(23)
```

**Feature engineering mattered  
more than the model family**

**Fuzzy matching with external CPU/GPU  
benchmark datasets → cpu\_mark, gpu\_mark**

**First-label simplification for multi-label  
categories**

**Vendor extraction (Intel/AMD/NVIDIA/Apple)**

**Frequency encoding of high-cardinality columns  
(Series, Brand)**

**One-hot encoding for low-cardinality categorical  
fields**

**Engineered performance ratios: gpu/cpu ratio,  
storage ratios**

# Feature Engineering & Encoding

Fuzzy matching with external CPU/GPU  
benchmark scores → cpu\_model

First-label categories

Vendor extension  
(Intel/AMD)

## Scaling

We selected only the continuous hardware-related variables for standardization:

- CPU/GPU benchmark scores, clock speeds, core/thread counts
- RAM size and frequency, SSD and HDD capacities
- Screen size, physical dimensions, weight, brightness, refresh rate
- Battery capacity, numeric counts (offers\_num, os\_bits)


These features have very different ranges and units, so scaling helps models that are sensitive to feature magnitude.

Feature engineering mattered more than the model family

# Feature Selection

- Used **Random Forest** on unscaled data to get 'feature\_importances'
- Kept variables with importance  $\geq 5\%$  of max  $\rightarrow 17$  features.
- Force-kept key specs: RAM, CPU/GPU benchmarks, SSD size, screen size.
- Added interaction features (total\_performance, total\_storage, perf\_per\_core, etc.)
- Grouped selected features into UI blocks (Memory, CPU, GPU, Display, Physical, Power) and ranked blocks by total importance

Block	Features
Memory & Storage	ram_gb, ssd_gb, hdd_gb, ram_freq_mhz
Processor	cpu_mark, cpu_cores, cpu_turbo_ghz
Graphics	gpu_mark, vram_gb, is_gaming
Display	screen_size_inch, refresh_rate_hz
Physical	height, width, weight_kg
Power	battery_wh, has_battery



**Result: A compact, domain-aware feature set that keeps the strongest price drivers.**



**Higher importance = more impact on price  
= should be collected first**

# Model Validation Strategy

## Validation Strategy

The approach ensured balanced representation of different price ranges across all splits, which is crucial given the skewed nature of the target variable. By reserving 20% of the data for final testing and using 60% for training and 20% for validation, we maintained a clear separation between tuning and evaluation phases. Stratified validation allowed for fair and reliable comparison across algorithms and highlights their real-world predictive strength.

## price\_num

target variable  
(numeric price)

A robust validation strategy by **splitting the cleaned dataset into training, validation, and test** sets using stratification based on price quantiles.

## Three Way Stratified Split

### Stratification by Price




### 80/20 Split (Train+Val vs Test)



### Within Train+Val: 75/25 Split

# Model Exploration

	Model	MAE	RMSE	R2_val	R2_cv_mean	R2_cv_std
1	GradientBoosting	342.649311	573.485327	0.716848	0.654554	0.033694
0	RandomForest	328.245928	615.404509	0.673941	0.702277	0.041281
3	XGBoost	325.667176	620.890166	0.668102	0.696091	0.064594
2	ExtraTrees	330.822773	654.513398	0.631182	0.667281	0.066521
8	KNN	369.904232	671.098478	0.612253	0.549843	0.047412
6	Lasso	447.919084	712.483101	0.562957	0.497915	0.051630
5	Ridge	447.942217	712.502052	0.562933	0.497910	0.051650
4	LinearRegression	447.966013	712.515468	0.562917	0.497904	0.051644
7	ElasticNet	447.255617	713.186523	0.562093	0.495328	0.051901
9	SVR	442.676372	807.562470	0.438529	0.347372	0.039810



Used both scaled and unscaled datasets depending on algorithm



Cross-validated on validation split (20%)

- Compared 10 models on the validation split using the same selected features:
- **Tree ensembles**  
(RandomForest, ExtraTrees, GradientBoosting, XGBoost\*) on unscaled data
- **Linear & distance-based**  
(Linear, Ridge, Lasso, ElasticNet, kNN, SVR) on scaled data

# Hyperparameter Tuning

Metric	Before	After	Change
<u>R<sup>2</sup></u>	0.640	0.714	+11.5%
<u>RMSE</u>	746	665	-81
<u>MAE</u>	426	357	-69


## GridSearchCV with 5-Fold Cross-Validation

- 200 trees with depth 5 → captures complex patterns without overfitting
- Learning rate 0.15 → aggressive but balanced by more trees
- Low regularization → data is clean, model benefits from flexibility

Optimal:

- n\_estimators=200
- max\_depth=5
- learning\_rate=0.15

# What Drives Computer Prices? (SHAP Analysis)




## Key Insights

1. **RAM has the highest impact (~€404 average)** - the dominant price driver
2. **Battery capacity (~€232)** - proxy for premium/ultrabook builds
3. **Number of offers** - market competition affects pricing
4. **CPU turbo speed & NVIDIA GPU**: performance specs matter

*A small set of features explains most of the price variation - RAM alone has 2x the impact of any other feature*

# How Features Affect Price (Non-Linear Relationships)



## Key Insights

- RAM: Linear with premium jump at high configs
- Battery: Interacts with form factor (thin + powerful = premium)
- Market competition shows diminishing marginal effects

# System Architecture & Inference Pipeline

# Offline Training

- We select the winning algorithm (Gradient Boosting) which performed best in our Notebook analysis.

# train\_model.py

- Purpose:  
Orchestrates the entire ETL pipeline.  
It cleans data,  
imputes missing values, and trains the model.

# Persistence Layer

We serialize the fitted model object (internal parameters) into a portable file  
`price_predictor.joblib` `feature_info.json`

Purpose: A binary file created by the Joblib library. It does not save the price; it saves the Trained Model Object (the 100 decision trees and their weights).

Purpose: The Schema Contract.  
Stores specific metadata like  
median\_ram=16GB and Apple=4.5  
needed for preprocessing.

UI

**The Streamlit App loads this saved model to execute inference on demand.**

Purpose: The Streamlit frontend.  
Captures raw user input  
(sliders/dropdowns).

# Live Market Analytics

## Browse Computers

### Smart Recommendations K-Nearest Neighbors

- Inside `app_utils.py`, we run `find_similar_products()`.
- Maps input to vector space, calculates Euclidean Distance.
- Returns 5 mathematical “twins” from 8,000+ database
- Why: This is purely mathematical search, requiring no trained model, just the raw data in memory.

## Market Segmentation (Analyst View)

### Market Segmentation K-Means Clustering

- The Logic: We run this Live (`perform_clustering`) on the DataFrame.
- Automatically detects "Families" of computers (e.g., Budget, Gaming, Workstation) based on Specs & Price.
- Because it runs live, we can visualize the market structure dynamically
- Offers business insights beyond simple price prediction.

# The Agentic Chatbot

## The Brain: Gemini 2.5 Flash with Function Calling (Tools)

- We integrated Google's Gemini 2.5 Flash model to provide a conversational interface.
- But purely generative models hallucinate prices. To fix this, we gave it "Hands".

### Tool 1: Price Prediction

```
get_price_prediction_for_agent:
```

LLM Action: "User asked for price." -> Tool Call: Runs our specific XGBoost .joblib model. -> Result: Returns precise € calculation.

### Tool 2: Recommendations

```
recommend_laptops_for_agent:
```

LLM Action: "User wants recommendations." -> Tool Call: Runs our KNN function. -> Result: Returns actual inventory.

### Tool 3: Data Query

```
get_brand_count:
```

LLM Action: "How many Dells?" -> Tool Call: Runs a Pandas count query. -> Result: 100% factual answer.

A Chatbot that calculates, not just chats

# Advancements, Improvements & Future Work

## Model

- Target  $R^2$  above 0.80 with hyperparameter tuning
- Ensemble stacking: combine GradientBoosting + XGBoost + RandomForest

## Data

- Expand dataset with more products and retailers
- Extract text features from descriptions using NLP

## Predictions

- Validate input combinations (realistic specs only)
- Recommend similar computers based on prediction

## UI

- Comparison mode: evaluate configurations side-by-side

# Advancements, Improvements & Future Work

## Uncertainty estimation & Prediction intervals

- Use Quantile Regression Forests or bootstrapped ensembles.
- Improve user trust by presenting ranges instead of point estimates.

## Device-Type Routing Models

- Separate models for desktops vs laptops vs ultra-portables.
- Each has distinct price dynamics → boosts accuracy noticeably.

## Hardware-Aware Neural Models

- Build a hybrid model combining engineered features with embeddings of hardware names (CPU/GPU text embeddings).

## Dynamic Market Adaptation

- Integrate periodic retraining schedules.
- Track distribution drift → alert when model becomes outdated.

## Similarity & Recommendation Engine

- kNN + cosine similarity for “best value for money”
- Bridge predictive → prescriptive functionality

## Improved Benchmark Matching via NLP models

- Replace regex matching with transformer embeddings.
- Increase CPU/GPU benchmark coverage significantly.

## Dockerised Deployment + CI Pipeline

- Full reproducibility.
- Auto-builds, versioning, and model registry.

## Limitations & Risks

- Dataset reflects only one marketplace → potential bias
- No temporal component → cannot capture price seasonality

# Thanks for Listening!



Group 6

